

Some Performance Optimizations for Embedded Systems

Dr. Jason Forsyth

Department of Engineering
James Madison University

Overview

- Moving from desktop/prototype to embedded implementation can be challenging; limited energy, data, and CPU cycles
- Calculations performed on different architecture may be significantly slower and/or infeasible on chosen platform.
- Several non-compiler arithmetic optimizations are available

Atmel 32u4

- Up to 16 MIPS; 16/32KB Program Memory; 1.25/2.5 KB RAM
- No floating point unit; no cache;
- 8 bit ALU but can calculate results of additional length via compiler unroll
- Native Data Sizes: int (2 bytes), long/float/double (4 bytes), long long (8 bytes)

Reducing Time Awake

- To achieve battery lifetimes must minimize the amount of time that the processor is awake.
- Major tasks: count steps, display steps, store data...etc.
- How long do each of these take? How can their efficiency be improved?
- Recall speeds: register < flash < SPI < i2c...etc.; int < float

Timing Your Code

- How long does a certain operation take? Use the internal cycle counter to find out.

Use micros() call

```
unsigned long time;

void setup(){
  Serial.begin(9600);
}

void loop(){
  Serial.print("Time: ");
  time = micros();

  Serial.println(time);
  delay(1000);
}
```

```
long startTime = micros(); ← Start Time Stamp
for (int i = 0; i < LOOPS; i++)
{
  //do something
}
long endTime = micros(); ← End Time Stamp

long delta = endTime - startTime;

Serial.print("Direct Timing: ");
Serial.println(delta);
```

Do it "many" times to get accurate timing

Ensure that whatever code is being timed can be "optimized away" by the compiler. Must do "something" with the final result. Default compile flag is `-Os` (which is `-O3` but for size)

Major Bottlenecks in Code

- Memory operations: read/writing to flash is slower than RAM; accessing any external I/O is significantly slower; excessive reads/writes when not necessary
- Arithmetic operations: excessive calculations that aren't necessary (values that never change); expensive operators;
- Algorithmic complexity: bad data structure; poor solution to problem

Measuring the Impact of float/int operations

```
#define LIMIT 256
#define TYPE float

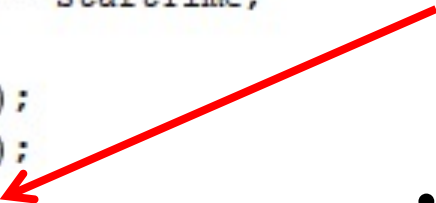
TYPE data[LIMIT];
TYPE result[LIMIT];

long startTime = micros();
TYPE sum = 0;
for (int i = 0; i < LIMIT; i++)
{
    sum = result[i] = (10 * data[i]);
}

long endTime = micros();

long delta = endTime - startTime;

Serial.print("End: ");
Serial.println(delta);
Serial.println(sum);
```



- Many digital filter (and desktop code) utilize floating point operations. However these are not naturally supported.
- Perform floating and integer calculation of arrays of various lengths and compare execution time.
- Code can do “dummy” work but must flush/use result so not removed by the compiler.
- Examine run time, program, and data size

Program and Data Size

```
#define LIMIT 256
#define TYPE float

TYPE data[LIMIT];
TYPE result[LIMIT];

long startTime = micros();
TYPE sum = 0;
for (int i = 0; i < LIMIT; i++)
{
    sum = result[i] = (10 * data[i]);
}

long endTime = micros();

long delta = endTime - startTime;

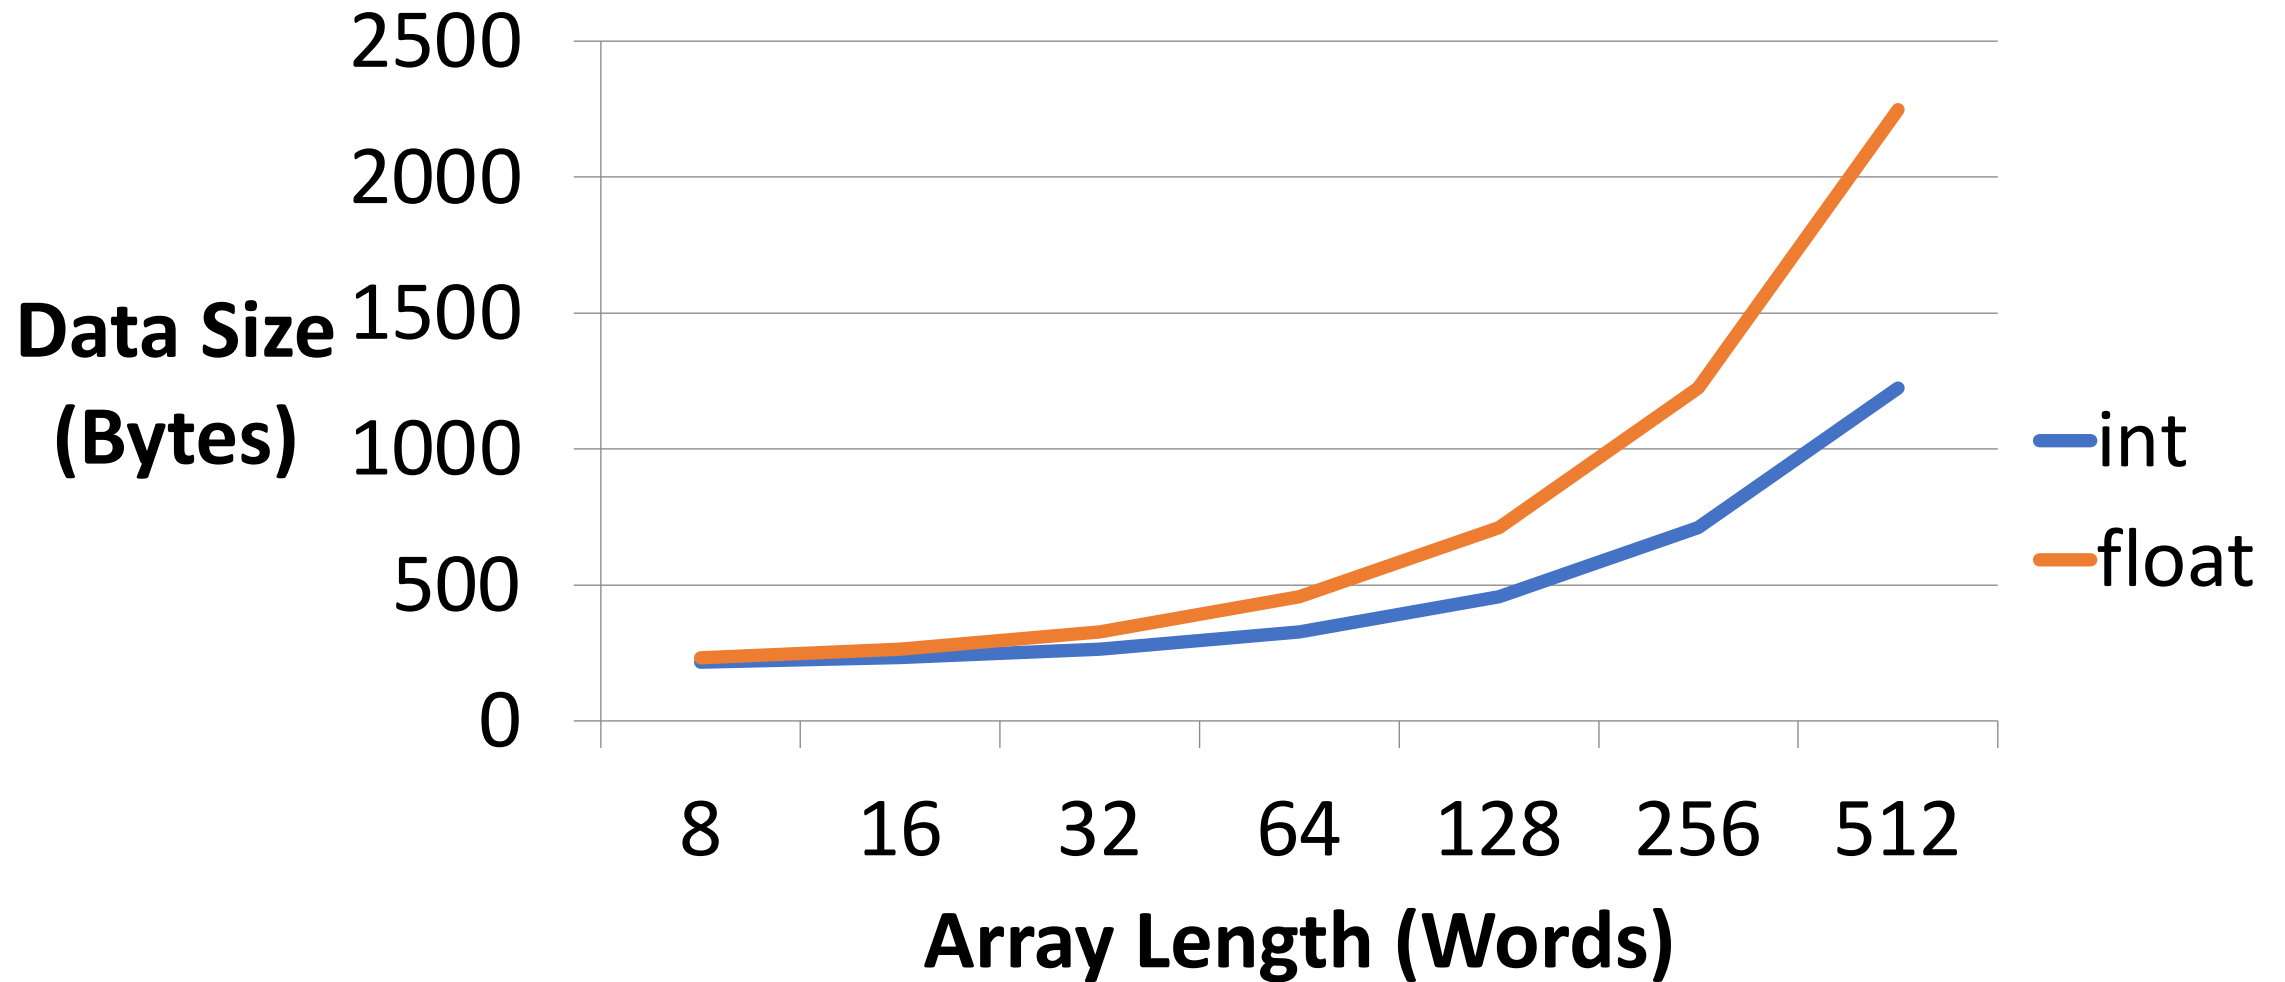
Serial.print("End: ");
Serial.println(delta);
Serial.println(sum);
```

		Array	Size		
Program Size (bytes)	8	16	32	64	128
int	1948	1948	1948	1948	1948
float	1950	1950	1950	1950	1950

Data Size (bytes)	8	16	32	64	128	256	512	1024
int	216	232	264	328	456	712	1224	2248
float	232	264	328	456	712	1224	2248	

- Program size is almost equal
- Data size grows faster (float is 2x bytes)
- Data Size = constant allocation + arrays

Program and Data Size



Execution Time

```
#define LIMIT 256
#define TYPE float

TYPE data[LIMIT];
TYPE result[LIMIT];

long startTime = micros();
TYPE sum = 0;
for (int i = 0; i < LIMIT; i++)
{
    sum = result[i] = (10 * data[i]);
}

long endTime = micros();


long delta = endTime - startTime;

Serial.print("End: ");
Serial.println(delta);
Serial.println(sum);
```

Execution Time (us)	8	16	32	64	128	256	512	1024
int	8	12	24	40	76	148	292	N/A
float	12	16	32	56	108	212	N/A	
Increase	50%	33%	33%	40%	42%	43%		

- Exec Time = constant stuff + loop time
- Floating point execution time is 30-50% slower
- Only doing 1 operation in this code. Would be worse if there were more...

Execution Time



So floats are bad... What to do?

- Know the cost and use them as little as possible.
- Move to a different processor? Is it worth the cost?
- “Integerize” calculations

“Integerizing” Calculations

- Floating point calcs can be turned into integers by moving the decimal point and then dividing the result
- Float: $y = 3.1435x + 3$
- Int = $y = (31435 * x) / 10000 + 3$
- Provides the “same” result but with potential losses due to integer multiplication and division

“Integerizing” Calculations

- Consider a digital filter: $y(x) = -2.386987234 * x[n] + 3.618872342 * x[n - 1] + 4.125872342 * x[n - 2]$
- Can be integerized with [-23870,36189,41259].
 - $y(x) = (-23870 * x[n] + 36189 * x[n - 1] + 41259 * x[n - 2])/10000$
- Cannot expand floats too many places as may result in overflow multiplication operation (n bits * n bits = 2*n bit result).
- Largest integer variable is *long long* at 8 bytes

“Integerizing” Calculations

- Decimal expansion is limited and will result in approximation errors
- Amount of acceptable error depends on application and filter structure (IIR may be problematic)

+/- 5E4			
X	Float calc	Int calc	Abs Error
15257	300010.9882	300013.1759	2.187719967
37431	240884.8791	240886.7805	1.90144279
48710	-17312.68037	-17312.512	0.168366942
37315	31688.13117	31688.4346	0.303433825
-8745	276615.662	276617.6136	1.951557713
36939	-26736.60544	-26736.5646	0.040836222

+/- 1E6			
X	Float calc	Int calc	Abs Error
705504	-1412453	-1412460	-7.03846
43352	4601.075	4601.341	0.265586
27798	107205.8	107206.7	0.814765
1814	289297.2	289299.3	2.082751
40475	89080.86	89081.71	0.849333
35666	163835.5	163836.8	1.260321

Accessing Memory

- Another significant delay to access memory that is "far away" or repeated access where not necessary
- Many filters and buffers may have data that is "shifted down" for calculation

```
for (int i = 0; i < LOOPS; i++)
{
    //move the data around to do the calculation
    sum = -2 * data[0] + data[1] + 4 * data[2];

    //shift the data down
    for(int j=0; j<LEN-1; j++)
    {
        data[j] = data[j+1];
    }
}
```


Circular versus Direct Buffers

- Two approaches: move the data (direct) or move the calculation around (circular)
- Can use modulus to adjust where calculation occurs. Always make relative to some header/pointer

```
for (int i = 0; i < LOOPS; i++)  
{  
    //move the calculation around on the data  
    sum = -2 * data[counter % LEN] + data[(counter + 1) % LEN] + 4 *  
    data[counter%LEN] = 0x37; //add some new data  
    counter++;  
}
```

Comparing Circular and Direct Buffer Implementations


```
TYPE sum=0;
long startTime = micros();
for (int i = 0; i < LOOPS; i++)
{
    //move the data around to do the calculation
    sum = -2 * data[0] + data[1] + 4 * data[2];


    //shift the data down
    for(int j=0;j<LEN-1;j++)
    {
        data[j] = data[j+1];
    }
}
long endTime = micros();

long delta = endTime - startTime;
```

```
int counter = 0;
sum = 0;
startTime = micros();
for (int i = 0; i < LOOPS; i++)
{
    //move the calculation around on the data
    sum = -2 * data[counter % LEN] + data[(counter + 1) % LEN]
        + 4 * data[(counter + 2) % LEN];
    data[counter%LEN] = 0x37; //add some new data
    counter++;
}
endTime = micros();

delta = endTime - startTime;
```




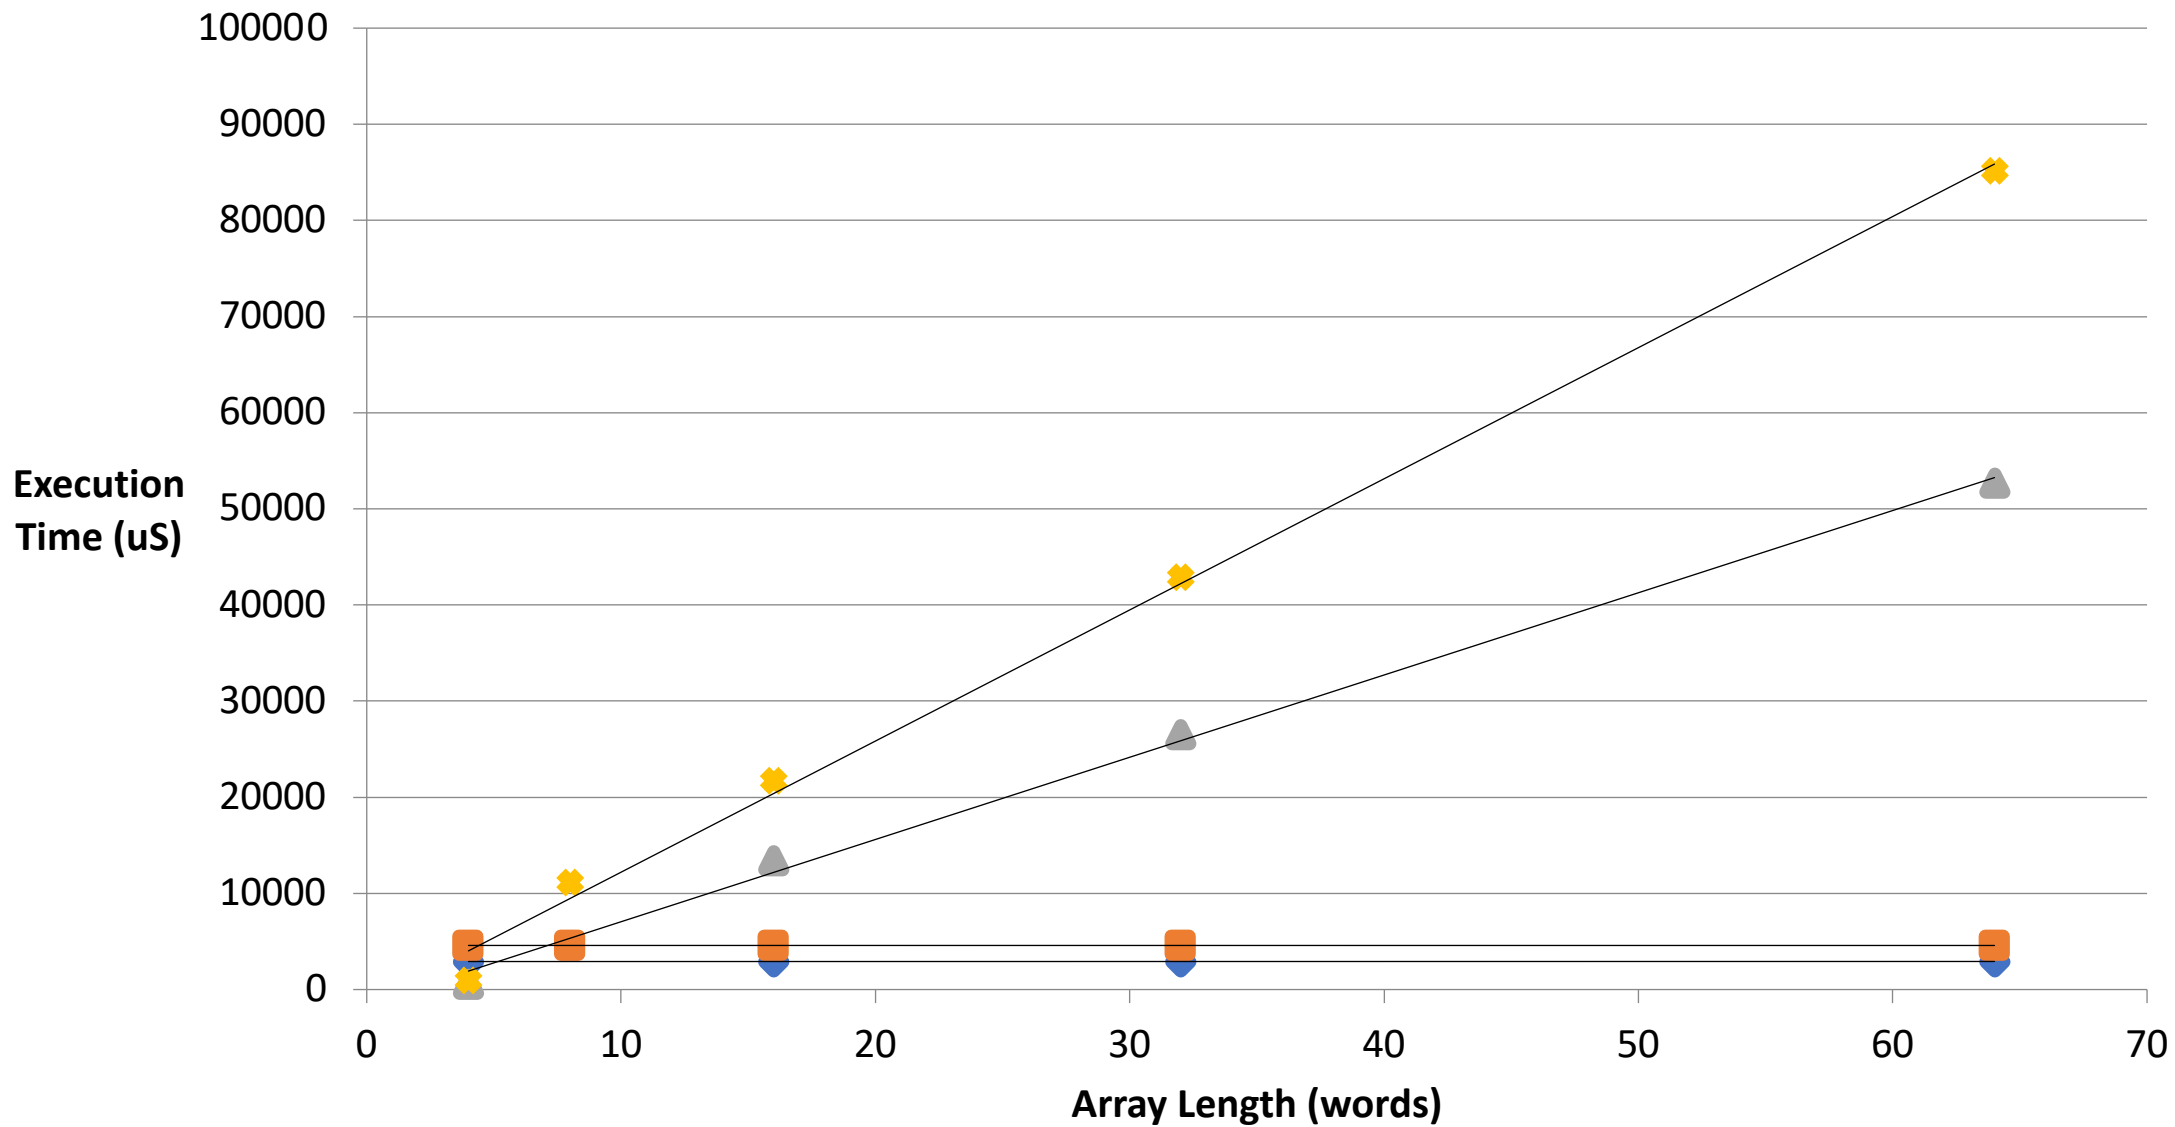


Comparing Circular and Direct Buffer Implementations

- In one slide direct is better, in the other circular is better...?
- Examine the array lengths. Cause significant difference between the algorithms.
- Modulus on 2^N value is a logical AND

**Execution Time on Various Length Integer
Array for Direct and Circular Calculations**





◆ Int Circular ■ float circular ▲ int direct ★ float direct
— Linear (Int Circular) — Linear (float circular) — Linear (int direct) — Linear (float direct)

Comparing Circular and Direct Buffer Implementations

- Calculation of modulus on not 2^N values “swamps” any benefits of the circular buffer
- Circular buffer is constant time operation $O(1)$
- Direct buffer is $O(N)$
- Float still more expensive than int operation.

How many optimizations to implement?

- Many (infinite) things to be done to make faster/better?
- Final requirement will be on step counting and current requirements
- $i_{avg} = 0.66 * i_{active} + 0.33 * i_{idle}$
- Reduction in running current is 2x more effective than reduction in sleep.
 - $\Delta i_{avg} = 0.66 * \Delta i_{active} + 0.33 * \Delta i_{idle}$

Other energy considerations

- Leave GPIO in “good” state. i2C pins left at 0V will drain current while asleep. Ensure all transactions are complete.
- Don't wake up too frequently. Use the watermark to delay running. Also, buffer EEPROM operations. Page write takes 4ms regardless.